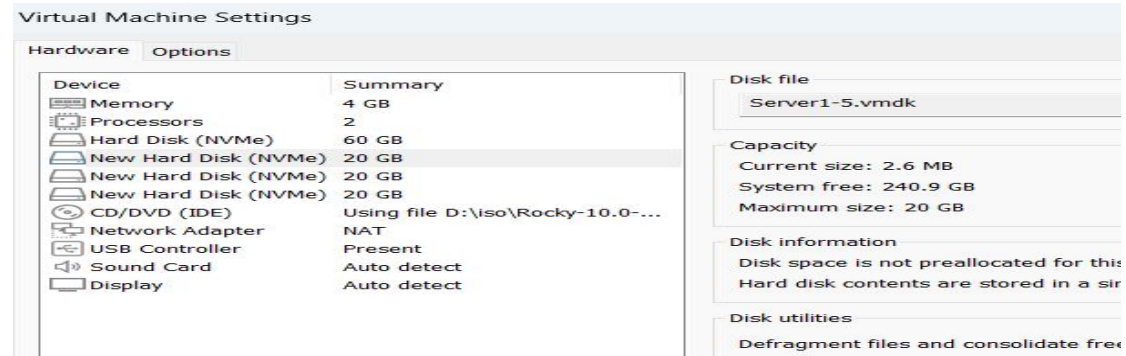


## Storage Management RAID

In this lab you will configure a RAID 1 and RAID 5 array in Linux.  
The mdadm package allows you to create RAID arrays.

First attach 3 hard disk of 20 GB each to the Linux virtual machine as shown below.



Then start the virtual machine and login with a user having sudo permissions.  
First you need to install the mdadm package.

```
sudo yum install mdadm -y
```

```
[admin@demosrv ~]# sudo yum install mdadm -y
```

You can verify the successful installation of mdadm package using the following command.

```
mdadm -h
```

```
[admin@demosrv ~]# mdadm -h
mdadm is used for building, managing, and monitoring
Linux md devices (aka RAID arrays)
Usage: mdadm --create device options...
        Create a new array from unused devices.
mdadm --assemble device options...
        Assemble a previously created array.
mdadm --build device options...
        Create or assemble an array without metadata.
mdadm --manage device options...
        make changes to an existing array.
mdadm --misc options... devices
        report on or modify various md related devices.
mdadm --grow options device
        resize/reshape an active array
mdadm --incremental device
        add/remove a device to/from an array as appropriate
mdadm --monitor options...
        Monitor one or more array for significant changes.
mdadm device options...
        Shorthand for --manage.
Any parameter that does not start with '-' is treated as a device name
or, for --examine-bitmap, a file name.
The first such name is often the name of an md device. Subsequent
names are often names of component devices.

For detailed help on the above major modes use --help after the mode
e.g.
mdadm --assemble --help
For general help on options use
mdadm --help-options
[admin@demosrv ~]#
```

Verify if there is any RAID array already created in the system using following command.

```
cat /proc/mdstat
```

```
[admin@demosrv ~]# cat /proc/mdstat
Personalities :
unused devices: <none>
[admin@demosrv ~]#
```

## Create a RAID 1 Array

First you will create a RAID 1 array. It is mirroring. It requires 2 hard disks. The RAID device is identified as /dev/md0 or /dev/md1 etc.

Identify the disk names for each 20 GB hard disk. Use lsblk command.

```
[admin@demosrv ~]$ lsblk
NAME                                MAJ:MIN RM  SIZE RO TYPE MOUNTPOINTS
sr0                                11:0    1   7.1G  0  rom
nvme0n2                           259:0    0   20G  0  disk
nvme0n3                           259:1    0   20G  0  disk
nvme0n4                           259:2    0   60G  0  disk
├─nvme0n4p1                       259:4    0    1M  0  part
├─nvme0n4p2                       259:5    0    1G  0  part /boot
├─nvme0n4p3                       259:6    0   59G  0  part
│   └─r1-root                     253:0    0   37G  0  lvm  /
│       └─r1-swap                 253:1    0   3.9G  0  lvm  [SWAP]
│           └─r1-home             253:2    0  18.1G  0  lvm  /home
└─nvme0n1                         259:3    0   20G  0  disk
[admin@demosrv ~]$ _
```

Here it is /dev/nvme0n1 , /dev/nvme0n2 and /dev/nvme0n3. In your case these names may be different.

Here you will use /dev/nvme0n1 and /dev/nvme0n2 hard disks to create a RAID array by name /dev/md0 .

To create a RAID 1 array give the following command.

```
sudo mdadm --create --verbose /dev/md0 --level=1 --raid-devices=2 /dev/nvme0n1 /dev/nvme0n2
```

```
[admin@demosrv ~]$ sudo mdadm --create --verbose /dev/md0 --level=1 --raid-devices=2 /dev/nvme0n1 /dev/nvme0n2
mdadm: Note: this array has metadata at the start and
may not be suitable as a boot device.  If you plan to
store '/boot' on this device please ensure that
your boot-loader understands md/v1.x metadata, or use
--metadata=0.90
mdadm: size set to 20954112K
Continue creating array [y/N]? y
mdadm: Defaulting to version 1.2 metadata
[ 1097.862123] md/raid1:md0: not clean -- starting background reconstruction
[ 1097.862183] md/raid1:md0: active with 2 out of 2 mirrors
[ 1097.863492] md0: detected capacity change from 0 to 41908224
mdadm: array /dev/md0 started.
[ 1097.867067] md: resync of RAID array md0
[admin@demosrv ~]$
```

Type y at the prompt to create the RAID array.

Verify that the RAID 1 array is created using following command.

```
cat /proc/mdstat
```

```
[admin@demosrv ~]$ cat /proc/mdstat
Personalities : [raid1]
md0 : active raid1 nvme0n2[1] nvme0n1[0]
      20954112 blocks super 1.2 [2/2] [UU]

unused devices: <none>
[admin@demosrv ~]$
```

Format the RAID array using following command.

**sudo mkfs.ext4 /dev/md0**

```
[admin@demosrv ~]$ sudo mkfs.ext4 /dev/md0
[sudo] password for admin:
mke2fs 1.47.1 (20-May-2024)
Creating filesystem with 5238528 4k blocks and 1310720 inodes
Filesystem UUID: cbc8c839-b374-4915-ad5d-8cd7e236d441
Superblock backups stored on blocks:
    32768, 98304, 163840, 229376, 294912, 819200, 884736, 1605632, 2654208,
    4096000

Allocating group tables: done
Writing inode tables: done
Creating journal (32768 blocks): done
Writing superblocks and filesystem accounting information: done

[admin@demosrv ~]$ _
```

Mount the array. For this create a directory by name /mnt/raid1.

```
[admin@demosrv ~]$ sudo mkdir /mnt/raid1
```

Now give following command to mount the /dev/md0 device.

**sudo mount -t auto /dev/md0 /mnt/raid1**

```
[admin@demosrv ~]$ sudo mount -t auto /dev/md0 /mnt/raid1
[ 1898.874800] EXT4-fs (md0): mounted filesystem cbc8c839-b374-4915-ad5d-8cd7e236d441 r/w with ordered dat
a mode. Quota mode: none.
[admin@demosrv ~]$ _
```

Now you can create data into /mnt/raid1 directory. These files will be written on both the hard disks.

To make sure that the array is reassembled automatically at boot, you have to adjust the /etc/mdadm/mdadm.conf file. You can automatically scan the active array and append the file with the following command.

**sudo mkdir /etc/mdadm**

**sudo mdadm --detail --scan | sudo tee -a /etc/mdadm/mdadm.conf**

```
[admin@demosrv ~]$ sudo mdadm --detail --scan | sudo tee -a /etc/mdadm/mdadm.conf
ARRAY /dev/md0 metadata=1.2 UUID=cff1eba0:480b4c6f:d35d67f7:8db247ef
```

You will need to add an entry in the /etc/fstab file for mounting the /dev/md0 device each time the Linux operating system starts.

This is how you have successfully created and mounted a RAID 1 array.

### **Removing the RAID 1 array.**

Unmount the /dev/md0 device using following command.

**sudo umount /dev/md0**

```
[admin@demosrv ~]$ sudo umount /dev/md0
[ 2160.214368] EXT4-fs (md0): unmounting filesystem cbc8c839-b374-4915-ad5d-8cd7e236d441.
[admin@demosrv ~]$ _
```

Then stop the RAID device /dev/md0

**sudo mdadm --stop /dev/md0**

```
[admin@demosrv ~]$ sudo mdadm --stop /dev/md0
[ 269.933584] md0: detected capacity change from 41908224 to 0
[ 269.933635] md: md0 stopped.
mdadm: stopped /dev/md0
[admin@demosrv ~]$
```

Verify with

**cat /proc/mdstat**

```
[admin@demosrv ~]$ cat /proc/mdstat
Personalities : [raid1]
unused devices: <none>
[admin@demosrv ~]$
```

Now give the following command

**lsblk -o NAME,SIZE,FSTYPE,TYPE,MOUNTPOINT**

```
[admin@demosrv ~]$ lsblk -o NAME,SIZE,FSTYPE,TYPE,MOUNTPOINT
NAME        SIZE FSTYPE  TYPE MOUNTPOINT
sr0         7.1G iso9660  rom
nvme0n1     60G                disk
├─nvme0n1p1  1M                part
├─nvme0n1p2  1G xfs           part /boot
├─nvme0n1p3  59G LVM2_member  part
│   ├─r1-root 37G xfs           lvm /
│   └─r1-swap 3.9G swap        lvm [SWAP]
└─r1-home 18.1G xfs           lvm /home
nvme0n2     20G linux_raid_member disk
nvme0n3     20G linux_raid_member disk
nvme0n4     20G                disk
[admin@demosrv ~]$ _
```

The output shows that the hard disks /dev/nvme0n2 and /dev/nvme0n3 are showing FSTYPE as linux\_raid\_member.

You need to zero their superblock which holds metadata for the RAID setup. Zeroing this removes the RAID metadata and resets them to normal. For this give following commands

**sudo mdadm --zero-superblock /dev/nvme0n2**  
**sudo mdadm --zero-superblock /dev/nvme0n3**

```
[admin@demosrv ~]$ sudo mdadm --zero-superblock /dev/nvme0n2
[admin@demosrv ~]$ sudo mdadm --zero-superblock /dev/nvme0n3
[admin@demosrv ~]$ _
```

Now again confirm with the lsblk command.

```
[admin@demosrv ~]$ lsblk -o NAME,SIZE,FSTYPE,TYPE,MOUNTPOINT
NAME        SIZE FSTYPE  TYPE MOUNTPOINT
sr0         7.1G iso9660  rom
nvme0n1     60G                disk
├─nvme0n1p1  1M                part
├─nvme0n1p2  1G xfs           part /boot
├─nvme0n1p3  59G LVM2_member  part
│   ├─r1-root 37G xfs           lvm /
│   └─r1-swap 3.9G swap        lvm [SWAP]
└─r1-home 18.1G xfs           lvm /home
nvme0n2     20G                disk
nvme0n3     20G                disk
nvme0n4     20G                disk
[admin@demosrv ~]$
```

Edit the /etc/mdadm.conf file and remove the RAID1 array entry.

Also if you have an entry in the /etc/fstab file for the /dev/md0 RAID 1 Array.

Restart the virtual machine using **sudo init 6**

## Create a RAID 5 Array

To create a RAID 5 array give the following command

```
sudo mdadm --create --verbose /dev/md0 --level=5 --raid-devices=3 /dev/nvme0n2 /dev/nvme0n3 /dev/nvme0n4
```

# make sure you give your system device names.

```
[admin@demosrv ~]# sudo mdadm --create --verbose /dev/md0 --level=5 --raid-devices=3 /dev/nvme0n2 /dev/nvme0n3 /dev/nvme0n4
[sudo] password for admin:
mdadm: layout defaults to left-symmetric
mdadm: layout defaults to left-symmetric
mdadm: chunk size defaults to 512K
mdadm: size set to 20954112K
mdadm: Defaulting to version 1.2 metadata
[ 2828.498116] raid6: skipped pq benchmark and selected avx2x4
[ 2828.498163] raid6: using avx2x2 recovery algorithm
[ 2828.512442] async_tx: api initialized (async)
[ 2828.522650] xor: automatically using best checksumming function   avx
[ 2828.576465] md/raid:md0: device nvme0n3 operational as raid disk 1
[ 2828.576528] md/raid:md0: device nvme0n2 operational as raid disk 0
[ 2828.582706] md/raid:md0: raid level 5 active with 2 out of 3 devices, algorithm 2
[ 2828.583209] md0: detected capacity change from 0 to 83816448
mdadm: array /dev/md0 started.
[ 2828.591886] md: recover of RAID array md0
[admin@demosrv ~]#
```

Check with

cat /proc/mdstat

```
[admin@demosrv ~]# cat /proc/mdstat
Personalities : [raid6] [raid5] [raid4]
md0 : active raid5 nvme0n4[1] nvme0n2[3] nvme0n3[0]
      41908224 blocks super 1.2 level 5, 512k chunk, algorithm 2 [3/3] [UUU]

unused devices: <none>
[admin@demosrv ~]#
```

Format the device using a file system. Here you will use EXT4 file system.

```
[admin@demosrv ~]# sudo mkfs.ext4 /dev/md0
[sudo] password for admin:
mke2fs 1.47.1 (20-May-2024)
/dev/md0 contains a ext4 file system
      last mounted on Thu Sep 25 14:24:20 2025
Proceed anyway? (y,N) y
Creating filesystem with 10477056 4k blocks and 2621440 inodes
Filesystem UUID: ef225a1e-24db-4598-9f4a-34dada86eceb
Superblock backups stored on blocks:
      32768, 98304, 163840, 229376, 294912, 819200, 884736, 1605632, 2654208,
      4096000, 7962624

Allocating group tables: done
Writing inode tables: done
Creating journal (65536 blocks): done
Writing superblocks and filesystem accounting information: done

[admin@demosrv ~]#
```

Now create a directory and mount the device.

Add an entry in the /etc/fstab file to mount the device each time the Linux system boots.

Follow the same steps as RAID 1 array to remove the RAID 5 array.